

# Load balancing problem with Least-Square Policy Iteration

Vincent Antaki

McGill University

# LSPI

## Least-Square Policy Iteration

Michail G. Lagoudakis & Ronald Parr [2003]. *Least-Square Policy Iteration*

# LSTDQ

```
LSTDQ ( $D, k, \phi, \gamma, \pi$ )           // Learns  $\hat{Q}^\pi$  from samples

//  $D$  : Source of samples  $(s, a, r, s')$ 
//  $k$  : Number of basis functions
//  $\phi$  : Basis functions
//  $\gamma$  : Discount factor
//  $\pi$  : Policy whose value function is sought

 $\tilde{\mathbf{A}} \leftarrow \mathbf{0}$            //  $(k \times k)$  matrix
 $\tilde{\mathbf{b}} \leftarrow \mathbf{0}$        //  $(k \times 1)$  vector

for each  $(s, a, r, s') \in D$ 
     $\tilde{\mathbf{A}} \leftarrow \tilde{\mathbf{A}} + \phi(s, a) \left( \phi(s, a) - \gamma \phi(s', \pi(s')) \right)^\top$ 
     $\tilde{\mathbf{b}} \leftarrow \tilde{\mathbf{b}} + \phi(s, a) r$ 

 $\tilde{\mathbf{w}}^\pi \leftarrow \tilde{\mathbf{A}}^{-1} \tilde{\mathbf{b}}$ 

return  $\tilde{\mathbf{w}}^\pi$ 
```

Figure: The LSTDQ algorithm

- ▶ Linear function approximation.
- ▶ Need to use basis functions.
- ▶ Requires pseudo-matrix inversions.

# Load balancing problem - variant 1

- ▶ The agent has 3 servers at its disposition and needs to dispatch them tasks it receives. Tasks arrive randomly following a Poisson distribution with  $\lambda = 2$ .
- ▶ Tasks requires a certain amount of work to be completed. This amount of work required to complete a task is equal to  $1 + T$  where  $T$  is a random poisson variable with  $\lambda = 5$ . The agent never knows the associated workload with a task.
- ▶ All server queues have a maximum length of 10. If the agent tries to add a task to an already full queue, the task is discarded and the agents receive a  $-5$  points reward.
- ▶ At every timestep, all servers accomplish one unit of work on the first task in their queue.
- ▶ Upon the completion a task by a server, the agent receives a reward equal to  $\frac{5}{\text{\# of iteration to complete task}}$ .

## Load balancing problem - variant 2

- ▶ We now have just 2 servers
- ▶ The amount of work generated by the servers is now different for every server and stochastic.
- ▶ Distributions :  $\mathcal{N}(0.8, 0.1), \mathcal{N}(1.2, 0.1)$

# Load balancing problem - state definition

For simplicity, we consider the state to be defined as follows :

- ▶ Current timestep
- ▶ Queue length of all servers
- ▶ Time since the current task has been added to the queue for all servers
- ▶ Number of timesteps spent on current task for all servers.

To notice :

- ▶ Multiples tasks can arrive at the same timestep.
- ▶ Decisions are not spread at a regular interval through time.

# Considerations

Difficulties inherent to that problem :

- ▶ Delayed reward
- ▶ A bit of stochasticity
- ▶ Partially observable state
- ▶ Would be better formulated as an average reward problem?

# Considerations

The optimal policy is however very simple.

- ▶ Give the task to the server which is expected to complete it the soonest.

For the first variant, this means :

- ▶ Give the task to the server with the shortest queue.
- ▶ If multiples servers have the shortest length of queue, give the task to the one which been running his task for the longest amount of time




# Choice of features - basis function dilemma

- ▶ Our  $q$ values are probably not a linear combination of the state as previously expressed. We need to consider basis functions to use for preprocessing.
- ▶ We do not wish to spend too much time on designing features; we would prefer the algorithm to sort it out.
- ▶ However, adding new features means we need to increase the amount of examples to look at (to avoid overfitting).
- ▶ Increasing the amount of examples is a costly option since LSTDQ does a pseudo-inverse operation.

# LSPI settings

- ▶ 100 episodes for policy evaluation
- ▶  $\gamma = 0.95$
- ▶ Our behavior policy is  $\epsilon$ -greedy with  $\epsilon = 0.2$
- ▶ Initially, we attempted  $\phi(s, a) = \text{concat}([\vec{0} \text{ if } a' \neq a \text{ else } s \forall a'])$  where  $\vec{0}$  is a null vector the same size as  $s$ . Results were unsuccessful.
- ▶  $\phi(s, a) : [\text{timestep}, l_{a'} - l_a \forall a' \neq a, t_{a'} - t_a \forall a' \neq a]$  where  $l_a$  is the length of the queue corresponding to action  $a$  and  $t_a$  the number of time since the first task of queue  $a$  was added to the queue.

# Results - variant 1



- Although these results are impressive, they can be mostly be credited to the use of the appropriate set of features.

## Results - variant 2

Unfortunately, our set of features did not achieve decent results on the second variant. Here are some features considered :

- ▶ the state
- ▶ the indicators  $l_{a'} = 0 \forall a'$  and, if  $l_{a'} \neq 0$ , ratios between queues length  $\frac{l_a}{l_{a'}} , \frac{l_{a'}}{l_a}$  else  $\vec{0}$ .
- ▶ Indicators over  $\frac{l_{a'}}{l_a} > r$  for  $r \in [0.1, 0.2, 0.25, 0.33, 0.5, 1, 2, 3, 4, 5, 10]$
- ▶ a onehot encoding of the action.
- ▶ Applying tanh, log out of despair.
- ▶ Action-based sparse encoding of input :  
 $\phi(x, a) = \text{concat}([\vec{0} \text{ if } a' \neq a \text{ else } x \forall a'])$
- ▶ For binary variables, we also considered the transformation  
 $\phi(x) = [x, 1 - x]$ .

In the end, none worked. We remain cursed with a linear model in a very non-linear environment.

# Conclusion

What have has been learn through those experiments?

- ▶ Methods like LSPI generate a new set of samples at each iteration and process it in one batch.
- ▶ Basis function dilemma.
- ▶ The load balancing problem (variant 2) is hard. Non-linear method should be considered.

# The End

Thank you!